

Paper2Repro: Execution-Grounded Paper to Code Reproduction

Anonymous ACL submission

Abstract

Automatic generation of executable code from scientific papers is crucial for both research aim and the practical adoption of new methods in industry. Existing paper-to-code agents can produce repositories that appear faithful and may score highly under LLM-as-a-judge evaluation, but such apparent quality does not ensure runnability or scientific reproducibility. We propose a rubric- and execution-guided paper-to-code framework that verifies generated repositories through intermediate-value tests, runtime checks, and benchmark-level reproduction criteria. Our workflow parses papers into structured JSON, generates initial repositories through a planning-analysis-coding pipeline, constructs paper-specific rubrics and `pytest` tests, extracts runtime requirements, and packages the resulting task into a Harbor-style evaluation environment. On 22 PaperBench papers, our prototype substantially improves runnability and reveals a large gap between executable repositories and scientifically reproducible repositories. These results suggest that paper-to-code evaluation should move beyond LLM-as-a-judge scores toward execution-grounded verification of algorithmic and experimental fidelity.

1 Introduction

Reproducibility is essential for scientific progress, allowing researchers to verify findings, build on previous work, and push the boundaries of knowledge (Open Science Collaboration, 2015; Baker, 2016; Pineau et al., 2021). In machine learning, however, reproducing experimental results from a paper alone remains a persistent challenge: documentation is often incomplete, hyperparameters are under-specified, and a substantial fraction of accepted papers release no accompanying code (Seo et al., 2025). This forces researchers to invest considerable effort reverse-engineering methods from prose, slowing the pace of scientific dissemination.

Recent Large Language Models (LLMs) have shown strong capabilities in both understanding scientific text and generating code (Dubey et al., 2024; OpenAI, 2024; Reid et al., 2024), prompting a wave of paper-to-code agents that aim to translate a research paper directly into an executable repository (Seo et al., 2025; Starace et al., 2025; Siegel et al., 2024a). Systems such as PaperCoder (Seo et al., 2025) decompose this task into planning, analysis, and coding stages, and emit repositories that are structurally complete and stylistically convincing. Under LLM-as-a-judge evaluation, such repositories routinely receive high faithfulness scores. The implicit conclusion is that the generated code is a reasonable proxy for the paper’s implementation.

We argue that this conclusion conflates three increasingly strict properties that should be kept distinct: whether a repository *looks* faithful to the paper, whether it actually *runs*, and whether it *reproduces* the paper’s experimental behavior. A simple analogy is egg fried rice. A finished plate may look right, but a positive visual judgment does not entail that the cook understood the recipe, followed the correct procedure, or produced a dish with the expected taste. A judge that only inspects the plate cannot detect a missing ingredient, an inverted step, or a heat profile that differs from the original. The same holds for paper-to-code systems: a repository may use the wrong loss term, omit a key normalization, depend on an uninstalled package, or execute successfully yet fail to match the reported baseline. Existing benchmarks largely measure surface-level repository quality and do not separate these three properties, leaving open the question of whether current paper-to-code agents are reproducing science or merely producing code that resembles it.

To address this gap, we propose that paper-to-code evaluation should be *execution-grounded* rather than *appearance-based*, and we present a prototype framework that realizes this principle. Given a paper, the framework runs three parallel

agents: a **coding agent** that emits an initial repository; a **rubric and unit-test generation agent** that converts paper claims into `pytest` cases targeting intermediate quantities and final outputs; and a **runtime-information extraction agent** that records hardware, package, dataset, and hyperparameter requirements. A subsequent code-running and amelioration stage uses rubric checks, test outcomes, and execution traces to revise the repository. The final artifact is packaged as a TerminalBench-style (Stanford and Anthropic, 2025) reproducibility task that can be replayed by external agents. The novelty here is not the multi-agent decomposition itself, but the shift in what counts as evidence of reproduction: rather than asking only whether the final result *looks* right, the framework verifies that the repository computes the correct intermediate quantities, runs to completion, and satisfies benchmark-level criteria.

We organize our investigation around three research questions:

RQ1 Does LLM-as-a-judge repository quality reflect actual runnability and scientific reproducibility?

RQ2 Can paper-specific rubrics and intermediate-value tests better identify algorithmic implementation errors than end-to-end output checks?

RQ3 Can static and runtime feedback meaningfully improve generated paper-to-code repositories?

We make three contributions. *First*, we identify and characterize a gap between apparent repository quality, executability, and scientific reproducibility in current paper-to-code agents. *Second*, we propose a rubric- and execution-guided workflow that decomposes paper-to-code reproduction into paper understanding, code generation, rubric and unit-test generation, runtime-information extraction, static amelioration, and TerminalBench-style task packaging. *Third*, we evaluate the prototype on 22 papers from PaperBench (Starace et al., 2025) and show that making generated repositories runnable is substantially easier than reproducing paper-level experimental claims, underscoring the need for intermediate-value verification and stronger runtime-guided refinement.

2 Related Work

Paper-to-code generation. PaperCoder (Seo et al., 2025) is the closest baseline to our work. It

decomposes paper-to-code into planning, analysis, and coding stages and shows that structured repository generation improves implementation quality over single-prompt baselines. However, PaperCoder’s main evaluation centers on repository-level faithfulness as judged by an LLM, with comparatively limited analysis of executability or experimental reproduction. Our framework adopts the planning-analysis-coding skeleton but extends it with paper-specific rubric generation, unit-test generation, runtime-information extraction, and a code-running amelioration loop intended to verify the generated repository against execution-grounded criteria.

Reproduction benchmarks. A growing body of work benchmarks paper-to-code systems and code-using agents. PaperBench (Starace et al., 2025) measures whether agents can reproduce machine-learning papers from scratch using hierarchical rubrics, and provides the papers we use in our evaluation. SciReplicate-Bench (Xiang et al., 2025) focuses on reproducing algorithmic procedures from method sections. CORE-Bench (Siegel et al., 2024b) evaluates the reproducibility of computational research more broadly, while REPRO-Bench (Hu et al., 2025) targets end-to-end reproduction with explicit execution checks. Compared to these benchmarks, our focus is not on building a static benchmark but on a *generative* workflow that turns a paper into a reproducible task on the fly, while inheriting their key insight that execution must be part of evaluation.

Agentic terminal benchmarks. Terminal-Bench (Stanford and Anthropic, 2025) evaluates agents on realistic terminal tasks using containerized environments, natural-language instructions, test suites, oracle solutions, and time limits. Our Harbor-style packaging adapts this evaluation philosophy to paper-to-code reproduction: the final artifact is not a static repository but a self-contained task that bundles paper, code, runtime specification, and tests.

Test-time refinement. Iterative self-debugging (Chen et al., 2023; Tian et al., 2024) and best-of- N sampling approaches (Brown et al., 2024) have shown that test-time feedback can improve LLM code generation. Our amelioration loop is in this lineage, but uses *paper-specific* rubrics and intermediate-value tests as the feedback signal, rather than generic correctness oracles or arbitrary

unit tests.

3 Method

In this section we describe our framework as a *stage-wise codebase transformation*. Rather than producing the final reproduction in one shot, the pipeline takes a paper R and walks the repository through a sequence of states $R_0 \rightarrow R_1 \rightarrow \dots \rightarrow R_n$, where each stage either *mutates the repository* or *emits a side artifact* that conditions later mutations. This view makes the framework easier to reason about and to debug: it tells us which stage *creates structure*, which stage *fills in behavior*, and which stage merely *packages* the result. The overall pipeline is shown in Figure 1.

The pipeline yields four artifacts: the generated code repository C , the paper-specific rubric together with an executable unit-test suite U , the structured runtime information I needed to execute the repository (datasets and environment), and a Harbor-style reproducibility task bundle B . The naive approach is to ask an LLM to generate C alone, as in prior repository-level code generation work (Seo et al., 2025); but C alone underspecifies reproducibility, because it conflates whether the code *runs* with whether it computes the *right* quantities. We therefore derive the evaluation artifacts U and I directly from R rather than from C , so that the paper—not the generated code—remains the source of ground truth.

We classify every stage by the kind of mutation it performs on the repository:

Create files that did not exist are written; repository *structure* appears.

Rewrite existing files are edited in place; *behavior* changes while structure stays stable.

Package/Spec the source is untouched; the stage emits a side artifact or wraps the repository.

A central observation of our analysis (Appendix B) is that the repository’s *structure* is established almost entirely by the early creation stages, whereas the later repair stage is overwhelmingly a *content rewrite* over a stable set of files. We describe each stage below following the build path and the evaluation path of Figure 1.

3.1 Plan

Objective. Produce the blueprint that conditions all later code generation, without writing any source code.

Mechanism. The cleaned paper is read once and distilled into an overall reproduction plan, an architecture decision (the set of files that will exist and their responsibilities), a configuration template, and a per-file logic analysis. We prepend a pre-processing step that converts the input PDF into a cleaned JSON representation R' (`paper.json`), with citation spans, bibliography entries, headers, and footers removed, so that every downstream stage reasons over the same canonical input. Crucially, every equation in the paper is extracted and recorded so that its literal presence can be checked after synthesis.

Effect on the codebase. None—this is a **Spec** stage. The repository is still empty; the outputs are planning side artifacts (plan, architecture, analyses, extracted equations) that become the authoritative context for synthesis. This stage and the analysis substep follow PaperCoder (Seo et al., 2025) closely; our contribution is the explicit extraction of equations and reported metrics, which later lets the rubric ground individual tests in specific paper quantities.

3.2 Codebase

Objective. Materialize the entire planned codebase as complete, runnable code.

Mechanism. Files are generated in dependency order, each conditioned on the plan, its own natural-language analysis, and the previously generated files. Every file in the planned task list is emitted as finished code rather than a TODO stub: the entry-point module, the per-component modules, a requirements file pinning package versions, and a configuration file. Paper equations are emitted *literally*—coefficients appear as numeric constants rather than free parameters—and a bounded self-verification step may regenerate a file at most once if a required equation is found missing.

Effect on the codebase. **Create** at maximum scale: this is the moment the repository comes into existence, and it already contains essentially the full structural footprint that persists to the end of the pipeline. We refer the reader to Seo et al. (2025) for full details of the planning, analysis, dependency-aware coding, and debugging substeps that this stage shares with PaperCoder.

3.3 Codebase with Dataset Runner

Objective. Make the synthesized code *executable against real data and rival methods*.




Figure 1: Overview of the Paper2Repro workflow as a stage-wise codebase transformation. From the input paper R , two paths proceed and later converge. The **build path** (top) materializes and repairs the repository: Paper $\rightarrow$ **Plan** $\rightarrow$ **Codebase** $\rightarrow$ **Iteratively-Fixed Codebase**. The **evaluation path** (left) derives the paper-grounded standard: from R we first *select a small, feasible baseline* to build an **Evaluation Plan**, which conditions the **Rubric**. The rubric is the hub of the workflow: it drives the in-place repair of the build path via *hard-coded rubric items*, it is operationalized into the **Unit Test** suite, and it supplies *dataset-loading information* that grows the **Codebase with dataset runner**. The dataset-runner codebase and the unit tests converge in **Codebase in Docker**, where a *runtime evaluation-and-fix* loop produces the portable, gradeable **Final Codebase**. Each stage either *creates* structure, *rewrites* existing files in place, or *packages* the repository without touching its source (Section 3).

Mechanism. First, an **evaluation plan** is produced: a catalog of the datasets and baseline methods the paper uses, each annotated with a feasibility decision (does it fit the hardware and storage budget?). Because faithfully running every baseline is rarely affordable, this stage deliberately *selects a small, feasible baseline* as the comparison target (the *Select Small Baseline* edge in Figure 1). Conditioned on that plan and on the rubric’s dataset-loading information, the repository is then grown with the supporting machinery—data-preparation scripts, metric code, and dataset/baseline configuration files—and the real datasets and rival repositories are fetched into an assets area. A set of **baseline-runner modules** is added behind a single dispatch table, so any rival can be invoked through one entry point. This stage also absorbs runtime setup: a structured artifact I records hardware requirements (e.g., GPU memory, CUDA version), package and Docker requirements, dataset endpoints and split specifications, training requirements, hyperparameters and seeds, and the metrics to be matched. Treating runtime setup as its own artifact—knowing that a paper needs CUDA 12 and a specific Torch version is not the same as knowing

how to implement the model—lets us construct, verify, and reuse the environment independently of the code.

Effect on the codebase. Create of the runtime and evaluation layer: new infrastructure files and a baselines package. Behavior here is intentionally incomplete—the runners are first wired as *stubs* that delegate to a bundled fallback and are smoke-tested, deferring faithful implementation to the repair stage.

3.4 Iteratively-Fixed Codebase

Objective. Bring the executable-but-incomplete code into faithful agreement with the paper.

Mechanism. The paper is first distilled into a **rubric**: a checklist of testable requirements broken into code-development, execution, and result-analysis items (“is this method actually implemented?”, “is this exact hyperparameter present?”). Crucially, the rubric is conditioned on the paper and the structured plan artifacts, *not* on the generated repository C and *not* on the grading harness. If the rubric were conditioned on C , the evaluation standard could be implicitly relaxed to fit whatever the coding agent happened to produce, and the paper

would no longer serve as ground truth. The static repair loop then walks the rubric against the code and checks code-development items that can be judged without execution—that the loss function contains a required penalty term, that the configuration file uses the correct hyperparameter, or that a named module exists at the expected location. When an item fails, the LLM proposes minimal search–replace patches and the loop iterates, keeping a numbered backup of every file it touches.

Effect on the codebase. Rewrite, structurally stable. The file tree does not grow (modulo backups); the change is concentrated *inside* existing files—for instance, rewriting a baseline stub into a faithful implementation, or aligning a configuration value to a paper-mandated constant. Many reproduction errors are visible in source code and are corrected cheaply here, before any training run is launched.

3.5 Codebase in Docker

Objective. Make the faithful codebase *gradeable* and *portable*.

Mechanism. The rubric is operationalized into an executable `pytest` suite synthesized in two tiers.

Intermediate-value tests. These check algorithmic internals—loss terms, diffusion coefficients, normalization invariants, tensor shapes, or values derived from specific equations in the paper. They are the diagnostically richest part of the framework: a system can have the right final number for the wrong reason, but it is much harder to silently match a battery of intermediate quantities.

Final-comparison tests. These run benchmark-level checks: does the generated method match or outperform the selected baseline on the dataset and metric described in the paper? Final-comparison tests answer the externally visible question, while intermediate-value tests answer the internally faithful question. This division is the core methodological contribution of our framework: rather than asking “does the final result look right?”, it asks “did the implementation compute the important internal quantities correctly, *and* did it also produce the right final result?”. Tests that reference a symbol the code does not actually expose are commented out and skipped rather than deleted, so a single hallucinated reference cannot break the whole run.

A planned *runtime evaluation-and-fix* loop closes the pipeline: it reads *I*, builds the environment, downloads datasets and baselines, runs

`reproduce.sh`, executes the `pytest` suite, reads `scores/scores.json`, feeds failing tests and logs back to the debugger, and patches and reruns until tests pass or an iteration cap is reached. The repository is then **packaged**, without source edits, into a self-contained container task: a `Dockerfile` baking in Python, `pytest`, and pinned dependencies; a task descriptor; a verifier entry script; and the repository plus its downloaded assets copied into the image.

Effect on the codebase. Create (the test suite) followed by **Package** (the source is frozen, not modified), yielding the portable Final Codebase. At the time of writing, the components of the runtime loop—dataset downloading, search–replace debugging, unit-test execution, and Harbor packaging—exist in isolation; the top-level orchestrator that sequences them autonomously is ongoing work.

3.6 Recurring Design Principles

Four invariants recur across stages and explain the observed behavior of the pipeline.

P1: Equation-literal synthesis. Equations are emitted as code with constants inlined, and their presence is verified, making the synthesized code checkable against the paper rather than merely plausible.

P2: Stub-then-implement baselines. Rivals are first wired as executable stubs behind a uniform interface, then faithfully implemented during repair. This separates *integration* (does the runner dispatch?) from *fidelity* (is the method correct?).

P3: Grader-independent, rubric-driven repair. Repair is conditioned on the paper-derived rubric, not the grading harness, so fixes target genuine paper requirements and generalize instead of overfitting a held-out grader.

P4: Non-destructive editing. Repairs rewrite in place with numbered backups, and hallucinated test references are skipped rather than deleted, so the pipeline never loses prior work—which is what keeps the file count stable through the repair stage (Appendix B).

4 Experimental Setup

We now describe the experimental setup and the metrics used to distinguish appearance-based from execution-grounded reproduction quality.

Papers. We evaluate the prototype on 22 papers drawn from PaperBench (Starace et al., 2025),

spanning a mix of deep-learning, optimization, and benchmark-style contributions. The papers were selected to cover a range of structural complexity and dataset requirements.

Backbone. The backbone model for all agents is GPT-5.2. Calls are routed through the same client and use a shared system-prompt template per stage; only the stage-specific instructions $\mathcal{T}_{\text{stage}}$ and the cleaned input R' differ across calls.

Environment. Each generated repository is packaged into a Harbor-style task with paper JSON, generated code, runtime information, `pytest` tests, and a weighted reward script. Tasks are run under a 30-minute timeout on a local machine with an RTX 4060.

Metrics. We report three metrics that together separate superficial faithfulness from execution-grounded reproduction: (i) *LLM Judge Score*, evaluated by GPT-5.2 against a set of generated rubrics; (ii) *Runnable Rate*, whether the repository can execute without fatal errors; and (iii) *Reproduction Score*, a weighted score over code-development, execution, and result-analysis tests; we do not separately report an intermediate-value pass rate, since it is subsumed by this score.

5 Results

We report preliminary results on a 22-paper subset of PaperBench together with BE-CBO, used as an in-depth case study. We treat these numbers as a sanity check on the framework rather than as a final evaluation; absolute values are likely to shift once the runtime run-and-fix loop is completed and dataset and baseline retrieval mature.

5.1 Generated Repo Details

To probe how closely generated repositories obey the paper that produced them, we compared three Paper2Repro outputs against their corresponding papers, rubrics, and official reference repositories: BaM, APT (Adaptive Pruning and Tuning), and BBOX-ADAPTER. The reference repositories are used as a behavioral calibration target, not as a strict file-template oracle. Scores are rubric-weighted implementation-faithfulness estimates, so they reflect both code coverage and evidence that the paper’s experiments were actually reproduced. Table 1 gives the headline numbers, and the paragraphs below discuss each repository in turn.

BaM: Batch and Match. The BaM repository is the strongest of the three in terms of core mathe-

matical logic. It implements BaM, ADVI, Score-VI, Fisher-VI, and GSM-style algorithms with full-covariance Gaussian variational families, synthetic target distributions, per-iteration metrics, and shared experiment infrastructure. This matches the paper’s central claim that BaM is a score-divergence-based black-box variational inference method with closed-form Gaussian-family updates. Most paper-obedience gaps lie in experimental detail rather than the algorithm skeleton: there is no committed evidence of the required long seeded runs, grid-search sweeps, final figures, or reproduced quantitative results. PosteriorDB support is mostly generic scaffolding—target-specific log-posterior and score functions for the paper’s models are absent or replaced by toy mechanisms—and the deep generative experiment is partial, with CIFAR loading, VAE-style code, posterior scoring, and reconstruction metrics in place but no trained decoder artifacts or completed reconstruction outputs. BaM therefore obeys the paper at the level of method structure, but not at the level of full experimental reproduction.

APT: Adaptive Pruning and Tuning. The APT repository implements a plausible adaptive-pruning system. It contains an APT-like adapter layer with frozen base weights, input/output masks, dynamic rank expansion, salience-density pruning, rank allocation, two-stage training, distillation scaffolding, evaluation utilities, and RoBERTa/T5 configurations that encode many paper hyperparameters, making the repository recognizable as an implementation of the paper’s high-level idea: adaptively tune salient parameters while pruning unimportant structure to improve training and inference efficiency. The main weaknesses are paper-critical details: baselines are local proxies rather than the specified Mask Tuning and CoFiPruning implementations, and several method details diverge from the paper or reference behavior, including effective adapter scaling, multi-head attention mask indexing, outlier salience, Equation-9-style adapter salience, layer mapping for distillation, and KL prediction distillation. The empirical side is especially weak: the only observed run is a toy tiny-random RoBERTa SST-2 run with no pruning, while the paper requires 60% sparsity experiments, relative accuracy analysis, efficiency claims, and ablations. APT therefore captures the broad engineering direction of the paper, but does not obey it closely on detailed mechanics or experimental evidence.

Paper	Score	Leaves	Implementation profile	Paper faithfulness
BaM	0.308	789	Strongest on core variational-inference algorithms; weak on executed experiments and real PosteriorDB/deep-generative reproduction.	Partially faithful: the main algorithmic logic is present, but paper-scale evidence is mostly missing.
APT	0.246	123	Coherent adaptive-pruning framework with masks, salience scoring, dynamic ranks, and training utilities; empirical evidence is minimal.	Weakly faithful: captures the broad APT idea, but diverges on several paper-critical details and lacks reported paper results.
BBOX-ADAPTER	0.290	279	Modular black-box adapter pipeline with pair scoring, online stores, NCE-style training, reranking, and beam search; configs and results are incomplete.	Partially faithful: architecture is recognizable, but training/evaluation details and reproduced tables are largely absent.

Table 1: Rubric-weighted summary of three generated repositories. “Score” is the weighted implementation-faithfulness estimate over the rubric leaves produced for that paper; “Leaves” is the number of leaf rubric items used to compute it.

BBOX-ADAPTER. The BBOX-ADAPTER repository contains a substantial adaptation pipeline: scalar question–answer pair scoring, ranking/NCE-style adapter training, online positive/negative example storage, single-step reranking, sentence-level full-step beam search, dataset preparation, and metric utilities. It also recognizes several important paper settings from Appendix H.2, including maximum generation length 512, temperature 1.0, learning rate 5×10^{-6} , batch size 64, 6000 training steps, AdamW, and beam size 3. Dataset manifests show materialized GSM8K, StrategyQA, TruthfulQA, and text-only ScienceQA splits. The repository nonetheless falls short of paper-faithful reproduction: there are no committed adapter checkpoints, training logs, evaluation summaries, or reproduced numbers for the paper’s main tables; the default configuration is not runnable as-is because config keys expected by the loader do not match the provided YAML; the ranking loss path has an indentation bug affecting `RankingNCELoss.forward`; and the Equation 3 score regularizer is effectively disabled by default because the α value is not resolved from the expected configuration location. Several paper components—MLM ablation, fine-tuning baselines, plug-and-play results, cost trends, and the Mixtral white-box extension—are stubs or planned aggregations rather than completed experiments.

Cross-paper interpretation. Across all three repositories, the dominant pattern is the same: generated code often captures the high-level method and implements plausible scaffolding, but falls short on paper-specific detail and experimental

Metric	P2C	P2R
Runnable Rate	0%	100%
LLM Judge (coding)	65%	85%
Intermediate-Value Score	–	72%
Baseline Outperformance	–	15%

Table 2: Paper2Code (P2C) vs. Paper2Repro (P2R) on 22 PaperBench papers. Intermediate-value and baseline metrics are not reported for P2C because none of its repositories execute end-to-end under our protocol.

proof. BaM is most convincing on core algorithmic logic; APT is most coherent as an engineering framework for pruning and tuning, yet shows the largest gap between method scaffold and paper-scale results; BBOX-ADAPTER has the broadest pipeline for a black-box LLM adaptation workflow, but suffers from configuration and loss-path issues that make reproducibility fragile. These comparisons argue that generated repositories should be treated as starting points for reproduction rather than as faithful reproductions of the papers themselves: useful for studying how much of a method an automated coding system can reconstruct, but requiring substantial human correction before they can support paper-level empirical claims.

5.2 Paper2Code vs. Paper2Repro

Table 2 compares Paper2Code (our re-implementation of the planning-analysis-coding baseline) with Paper2Repro (our framework) on four metrics: runnable rate, LLM judge score on the coding stage, intermediate-value pass rate, and baseline-outperformance rate.

The most striking row is runnability. Under our evaluation protocol, none of the 22 P2C repos-

ries run end-to-end within the 30-minute timeout, while all 22 P2R repositories do. Two factors drive this gap: the runtime-information agent fixes most environment-level mismatches before execution, and the static amelioration loop catches the structural bugs—missing imports, mis-spelled module paths, mis-typed configuration keys—that otherwise prevent a repository from starting. Because no P2C repository runs, intermediate-value and baseline metrics cannot be evaluated for P2C, and the corresponding cells are left empty.

The second row shows that P2R also improves the LLM judge score on the coding stage from 65% to 85%. This is consistent with the original Paper2Code finding that decomposed generation produces faithful-looking repositories, and shows that adding amelioration does not regress on appearance-based quality.

The remaining rows tell a more cautious story. P2R passes roughly 72% of intermediate-value tests but matches or beats stated baselines on only about 15% of papers. In other words, even after the framework has produced a runnable repository whose individual algorithmic components mostly look correct, the overall system still rarely reproduces the paper’s reported numbers. Two factors contribute. First, part of the intermediate-value pass rate is achieved at the cost of correctness: when a generated step fails, the amelioration loop sometimes patches the failure by simplifying or removing the offending computation, which keeps the code runnable but lowers faithfulness on tests it can no longer reach. Second, baseline-outperformance is sensitive to factors outside the generated code itself, including dataset retrieval, baseline availability, and training budget—issues that the static amelioration loop is not designed to fix and that the runtime run-and-fix loop is intended to address.

5.3 Stage-wise Reproduction Scores

Figure 2 tracks the reproduction score as a repository walks the build path of Figure 1, from the Paper2Code baseline through synthesis (**Codebase**), static repair (**Iteratively-Fixed Codebase**), and baseline wiring (**Codebase with Dataset Runner**). The headline is that every stage clears the Paper2Code baseline, but the climb is not steady: the score can *fall* after the Iteratively-Fixed Codebase stage and after the Codebase with Dataset Runner stage rather than rising at every step.

Two mechanisms explain these dips. *First, real*

regressions. Both repair and baseline-wiring are search–replace edits over code that already runs, so satisfying one rubric item can break another that previously passed. The repair stage, for example, reliably broke *all-in-one*’s Hodgkin–Huxley noisy-time-series summary statistic—a leaf that passed right after synthesis and failed once repair had touched the surrounding file. The baseline-wiring stage is even clearer: in *bridging-data-gaps* it reliably broke the **adaptor module** and the **CDC baseline**—the very rival methods it is meant to wire in—and in *stochastic-interpolants* it broke the in-painting time-derivative interpolant. Because each edit is local but the rubric is global, a stage can add the feature it targets and silently disturb a neighbor. *Second, a nondeterministic grader*. The judge scores each rubric leaf by first asking an LLM to pick the relevant files from the repository, then reading only those and scoring 0/1 with an instruction to “treat missing as failure.” That file-selection step varies run to run, so a leaf can flip $0 \leftrightarrow 1$ on *identical* code. A single-run bar therefore mixes genuine movement with this flicker, which is why the cross-stage analysis trusts how many of three runs a leaf passes rather than any one run.

Read through the multi-run signal, the net picture is steadier than the single-run bars suggest. The Iteratively-Fixed Codebase stage is the consistent within-tool win: it contributes reliable (3/3) gains in all three papers—for instance `random_fourier_features=False` in *stochastic-interpolants*, the gradient-of-output-layer and 10-shot generation code in *bridging-data-gaps*, and the provided training loop plus the Webb-2018 edge-update algorithm in *all-in-one*. The Codebase with Dataset Runner stage is paper-dependent: it helps *stochastic-interpolants* but net-regresses *bridging-data-gaps*, precisely because it breaks the adaptor and CDC baselines it is supposed to add. The practical takeaway is to trust the multi-run stability of a stage rather than an isolated single-run dip.

5.4 Case Study: The Code Quality of BE-CBO Repo

BE-CBO is a useful case study because its acquisition function and constraint-handling logic admit crisp, paper-specific intermediate checks. Table 3 compares the components implemented in the paper’s reference code (BE-CBO-std), in the Paper2Code-only output (P2C), and in our framework’s output (P2R), on five components that




Figure 2: Reproduction score by pipeline stage for three papers (*bridging-data-gaps*, *stochastic-interpolants*, *all-in-one*). The four bars per paper follow the build path of Figure 1: the **Paper2Code** baseline, the freshly synthesized **Codebase**, the **Iteratively-Fixed Codebase** (static repair), and the **Codebase with Dataset Runner** (baseline wiring and dataset-loading). Every stage clears the Paper2Code baseline, but the score is *not monotone* across stages: a repair or baseline-wiring edit can lower the score relative to the stage before it.

intermediate-value tests directly verify.

Across all five rows, P2C diverges from the paper in ways that are visible to a careful reader but not to an LLM-as-a-judge: it substitutes a plain neural ensemble for a variational probabilistic ensemble, replaces the BoTorch constrained acquisition with a hand-rolled Sobol+Adam+SLSQP loop, and continues training the constraint classifier across BO iterations rather than retraining at each fit. P2R recovers the paper’s specified components on every row, because the rubric agent generates intermediate-value tests that pin the relevant choices in place, and the amelioration loop produces patches that satisfy them.

5.5 Per-Paper Score Distribution

Figures 3 and 4 show intermediate-value and final reproduction scores for each paper in the evaluation set, color-coded by difficulty tier.

Intermediate-value scores cluster between 70% and 85% for the majority of papers, with a mean of 72.8% (Figure 3). The top of the distribution is dominated by T5 “easier” papers—BE-CBO at 90%, stay-on-topic-with-classifier-free-guidance at 87%, semantic-self-consistency at 84%—and the bottom by papers in T0–T2 tiers, including a deliberately included fake-no-asset-paper con-

trol at 0% and one generation failure (sapg, 5%). The mostly flat distribution above 60% indicates that the framework can identify and implement the core algorithmic components for most papers in the set.

Reproduction scores fall in a substantially lower range, with a mean of 16.2% (Figure 4). BE-CBO sits at 45% as a T6 outlier with a custom rubric and task; the next tier (stay-on-topic-with-classifier-free-guidance 35%, semantic-self-consistency 30%) and the broad middle of T4 “medium” papers (mostly 20–22%) confirm the component-vs.-benchmark gap. T1 papers, which combine reinforcement learning with simulator dependencies, bottom out at 5–8%, reflecting how strongly final scores depend on factors outside the generated code.

The component-to-benchmark gap. The headline finding is the size of the gap between Figure 3 and Figure 4: roughly 56 percentage points on average. Even after a framework succeeds at making code runnable and at implementing most of the right algorithmic components, the residual challenge of matching benchmark-level numbers remains large. Diagnosing why this gap is so wide is

Component	BE-CBO-std (paper)	P2C	P2R (ours)
Constraint classifier	Variational probabilistic Deep Ensemble	Plain neural ensemble; hand-crafted probit loss	Variational probabilistic Deep Ensemble
Classifier training	Retrain at each fit	Continue training across BO iterations	Reinitialize before each refit
Objective surrogate	Explicit GP with stated modeling choices	Default GP wrapper	Explicit GP with stated modeling choices
Feasibility bound	Uncertainty-aware bound	Approximate $p + \sigma \geq 0.5$	Uncertainty-aware bound
Acquisition optimization	BoTorch constrained acquisition	Manual Sobol + Adam + SLSQP	BoTorch <code>optimize_acqf</code> with Sobol restarts

Table 3: BE-CBO implementation differences. P2R recovers the paper’s specified components on every row where P2C diverges; the rubric agent generates intermediate-value tests that pin each choice down, and the amelioration loop produces patches that satisfy them.

the central analysis remaining for the final version of this paper. *[Reason for the gap: TBD.]*

5.6 Replication Results of the Final Repo

The scores reported above are rubric-weighted, LLM-graded estimates of *implementation faithfulness*; they do not by themselves show a generated repository *executing* on real model weights. As a preliminary check that the final artifact actually runs, we executed one generated repository end-to-end—*bridging-data-gaps* (DPMs-ANT)—rather than only scoring it, giving a concrete instance of the runtime loop the rest of this paper treats as future work. The run used an RTX 4060 (8 GB), loading the real pretrained DDPM `google/ddpm-ema-celebahq-256` (an FFHQ-like face source standing in for the paper’s source domain) and a real 10-image target (`huggan/anime-faces`, the closest public stand-in for the paper’s curated 10-shot sets, which are not downloadable). We stress at the outset that these are **real numbers on real weights, but they are not the paper’s numbers**: they answer “does the adaptation pipeline run on real weights?” (yes), not “does it match the paper’s tables?”.

What moved. On the 10 target images, training loss dropped $0.0177 \rightarrow 0.0080$, and the held-out ϵ -MSE improved $0.00968 \rightarrow 0.00876$ ($\sim 9.5\%$; Table 4), i.e. the model genuinely shifted toward the target domain. The repo’s own Intra-LPIPS metric produced a real number (0.5101, over 12 generated samples at 50 DDIM steps), and the whole run took 86 s at a peak of 6.43 GB VRAM. The held-out gain is modest *by design*: only 60 adaptation steps, only 56 K norm-only parameters, and an evaluation averaged over all noise levels (high-noise timesteps are nearly domain-independent and dilute the signal). The training-loss drop is the clearer evidence

Model (eval-plan in-scope)	Held-out ϵ -MSE $\downarrow$
Traditional DDPM (frozen)	0.00968
Param.-efficient adaptation	0.00876

Table 4: End-to-end replication of the final *bridging-data-gaps* repository. Both rows load the same real pretrained DDPM (`ddpm-ema-celebahq-256`) and are evaluated on the same real 10-image target. “Traditional DDPM (frozen)” is the pretrained model with no adaptation—the honest “before” anchor; “Param.-efficient adaptation” fine-tunes only the 56 K GroupNorm parameters on the 10 target images, standing in for the paper’s “adaptor-only / few extra parameters” setting. Lower is better.

that real adaptation is happening.

What this run does and does not show. This is a deliberately reduced-scale demonstration, and several gaps separate it from the paper’s reported results. The target domain is anime faces rather than the paper’s curated sunglasses/babies/sketches; we run 60 steps rather than 300; we adapt by norm-tuning rather than with the ANT adaptor; there is no classifier or adversarial-noise stage; and Intra-LPIPS is computed over 12 samples rather than 1000. The ANT method proper was *not* run: its `train_classifier` $\rightarrow$ `train_ant` $\rightarrow$ `eval` pipeline enforces batch 40 / 300 iterations / 1000-sample evaluation and needs source images plus a classifier checkpoint, but batch-40 at 256 px does not fit 8 GB and a 1000-sample evaluation takes hours. We also hit a genuine codebase bug: the repo’s own `AdapterInjector` mis-wires modules on a real `diffusers` UNet (a 512-channel bottleneck block ends up executing at a 128-channel/256 px position, raising a GroupNorm channel mismatch), so to stay within budget we disabled the broken adapter and adapted via norm-tuning while still reusing the repo’s real




Figure 3: Estimated intermediate-value correctness per paper (prior, calibrated to a 72.8% mean). Difficulty tiers run from T0 (generation failed) to T5 (easier, LLM-only).

`DiffusionModelWrapper.predict_eps`, `NoiseSchedule`, and `Intra-LPIS` core. What the run *does* establish is concrete: real pretrained weights load through the repository’s wrapper, the real forward-noising $\rightarrow \epsilon$ -prediction $\rightarrow$ adaptation loop runs end-to-end on real data within 8 GB, the loss moves the right way, and the repository’s metric yields a real number—the gap to matching the paper’s tables is one of compute budget and access to the paper’s data, not of a pipeline that fails to run.

6 Discussion

Answering RQ1. The gap between LLM judge scores and intermediate-value pass rates supports the claim that “seems good” is not enough. A judge that inspects only the final repository can miss a missing penalty term, an off-by-one normalization, or an incorrect optimizer setting that does not prevent the code from running.

Answering RQ2. Intermediate-value tests are the strongest diagnostic component of the framework. In several case studies, an intermediate-value test caught an error that neither LLM judge inspection

nor final metric comparison could isolate, because the final metric was either similar to baseline by coincidence or unreachable due to a separate runtime issue.

Answering RQ3. The static amelioration loop already provides substantial improvements in code-development quality at low cost. The full answer to RQ3, however, requires the completed runtime run-and-fix loop, which we leave to future work.

Limitations

This work has several limitations. First, the runtime run-and-fix loop is not fully implemented; while the components exist, the top-level orchestrator that ties them into a single autonomous process is future work. Second, our experiments cover 22 papers under a single backbone (GPT-5.2) and a single hardware budget; broader coverage and multiple-seed runs are needed before any conclusions about absolute performance can be drawn. Third, our intermediate-value tests are themselves LLM-generated from the paper; while we keep them blind to the generated repository, they may still inherit biases from the backbone, and a care-




Figure 4: Estimated final reproduction score per paper (outperforms stated baseline; prior). Mean is 16.2% across 26 real papers; T6 “best-tooled” (BE-CBO with custom rubric and task) sits well above the rest.

ful audit of test quality is needed. Fourth, dataset and baseline fetching are not always reliable, and some failures are infrastructure failures rather than implementation failures. Fifth, we do not yet run formal ablations isolating the contribution of each stage.

7 Future Work

Full runtime orchestrator. The most immediate next step is to implement the full runtime run-and-fix orchestrator that sequences environment build, dataset retrieval, `reproduce.sh` execution, test running, log ingestion, and patch application into a single loop bounded by an iteration cap.

Stronger environment builder. Many papers require specific CUDA versions, Torch versions, GPU memory budgets, or non-Python dependencies. A better environment builder should translate the runtime-info JSON into a concrete `Dockerfile` and installation script, with verifiable post-build checks.

Improved dataset and baseline fetching. Preliminary results suggest that low reproduction scores are often caused by dataset retrieval or base-

line availability, not by model implementation errors. Targeted improvements in this area are likely to raise reproduction scores more than additional amelioration of model code.

Formal ablations. Important ablations include removing intermediate-value tests, removing static amelioration, removing runtime feedback, and comparing against a PaperRepro-only baseline under matched conditions. These ablations would quantify the contribution of each component of the framework.

Stronger case studies. BE-CBO illustrates how intermediate-value tests detect whether key algorithmic components are correctly implemented. Negative cases are equally informative, since they show how repositories may be runnable yet scientifically incorrect, and they argue most strongly for execution-grounded evaluation.

8 Conclusion

PaperRepro asks whether an LLM can generate a faithful-looking repository from a paper. Terminal-Bench asks whether an agent can solve realistic terminal tasks under tests. Our project sits be-

tween the two and asks whether a paper-to-code system can generate a repository that not only looks good, but also runs and satisfies paper-specific scientific reproduction tests. We present a rubric- and execution-guided framework that turns a paper into an initial repository, paper-specific rubrics, intermediate-value tests, runtime requirements, and a Harbor-style reproducibility task. On 22 PaperBench papers, the prototype substantially improves runnability and exposes a large gap between executable and scientifically reproducible repositories. We view this as evidence that paper-to-code evaluation should move beyond LLM-as-a-judge scores toward execution-grounded verification of algorithmic and experimental fidelity.

References

Monya Baker. 2016. [1,500 scientists lift the lid on reproducibility](#). *Nature*, 533(7604):452–454.

Bradley Brown, Jordan Juravsky, Ryan Ehrlich, Ronald Clark, Quoc V. Le, Christopher Ré, and Azalia Mirhoseini. 2024. [Large language monkeys: Scaling inference compute with repeated sampling](#). *arXiv preprint arXiv:2407.21787*.

Xinyun Chen, Maxwell Lin, Nathanael Schärli, and Denny Zhou. 2023. [Teaching large language models to self-debug](#). *arXiv preprint arXiv:2304.05128*.

Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, and 1 others. 2024. [The Llama 3 herd of models](#). *arXiv preprint arXiv:2407.21783*.

Chuxuan Hu, Liyun Zhang, Yeji Lim, Aum Wadhvani, Austin Peters, and Daniel Kang. 2025. [Repro-bench: Can agentic ai systems assess the reproducibility of social science research?](#) In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 23616–23626. Association for Computational Linguistics.

Open Science Collaboration. 2015. [Estimating the reproducibility of psychological science](#). *Science*, 349(6251):aac4716.

OpenAI. 2024. [GPT-4 technical report](#). *arXiv preprint arXiv:2303.08774*.

Joelle Pineau, Philippe Vincent-Lamarre, Koustuv Sinha, Vincent Larivière, Alina Beygelzimer, Florence d’Alche Buc, Emily Fox, and Hugo Larochelle. 2021. [Improving reproducibility in machine learning research: A report from the NeurIPS 2019 reproducibility program](#). *Journal of Machine Learning Research*, 22(164):1–20.

Machel Reid, Nikolay Savinov, Denis Teplyashin, Dmitry Lepikhin, Timothy P. Lillicrap, Jean-Baptiste Alayrac, Radu Soricut, Angeliki Lazaridou, Orhan Firat, Julian Schrittwieser, and 1 others. 2024. [Gemini 1.5: Unlocking multimodal understanding across millions of tokens of context](#). *arXiv preprint arXiv:2403.05530*.

Minju Seo, Jinheon Baek, Seongyun Lee, and Sung Ju Hwang. 2025. [Paper2code: Automating code generation from scientific papers in machine learning](#). *arXiv preprint arXiv:2504.17192*.

Zachary S. Siegel, Sayash Kapoor, Nitya Nadgir, Benedikt Stroebl, and Arvind Narayanan. 2024a. [Core-bench: Fostering the credibility of published research through a computational reproducibility agent benchmark](#). *Transactions on Machine Learning Research*.

Zachary S Siegel, Sayash Kapoor, Nitya Nadgir, Benedikt Stroebl, and Arvind Narayanan. 2024b. [Core-bench: Fostering the credibility of published research through a computational reproducibility agent benchmark](#). *arXiv preprint arXiv:2409.11363*.

Stanford and Anthropic. 2025. [Terminal-bench: A benchmark for ai agents in terminal environments](#). <https://www.tbench.ai/>. TODO: replace with canonical citation when available.

Giulio Starace, Oliver Jaffe, Dane Sherburn, James Aung, Jun Shern Chan, Leon Maksin, Rachel Dias, Evan Mays, Benjamin Kinsella, Wyatt Thompson, Johannes Heidecke, Amelia Glaese, and Tejal Patwardhan. 2025. [Paperbench: Evaluating ai’s ability to replicate ai research](#). *arXiv preprint arXiv:2504.01848*.

Runchu Tian, Yining Ye, Yujia Qin, Xin Cong, Yankai Lin, Yinxu Pan, Yesai Wu, Haotian Hui, Weichuan Liu, Zhiyuan Liu, and Maosong Sun. 2024. [Debugbench: Evaluating debugging capability of large language models](#). *arXiv preprint arXiv:2401.04621*.

Yanzheng Xiang, Hanqi Yan, Shuyin Ouyang, Lin Gui, and Yulan He. 2025. [Sci replicate-bench: Benchmarking llms in agent-driven algorithmic reproduction from research papers](#). *arXiv preprint arXiv:2504.00255*.

A Workflow Diagram

The high-level workflow diagram is presented in Figure 1 in the main text. It depicts the five stages described in Section 3, the data dependencies between them, and the placement of the static repair loop and the planned runtime evaluation-and-fix loop.

B Case Study: *all-in-one* Codebase Evolution

To ground the stage-wise transformation of Section 3 in real numbers, we trace the reproduction of

Metric	Stage 3	Stage 8.1	Δ
Source files	164	164	0
Python LOC	51,049	51,223	+174
Files edited in place	—	15	15
Backups (.bak) added	0	16	+16

Table 5: Stage 3 (post-synthesis) $\rightarrow$ Stage 8.1 (post-repair) repository delta for *all-in-one*. The entire source delta is +415/ $-$ 163 lines (net +252) across 15 edited files; the file count is unchanged.

“All-in-one simulation-based inference” (hereafter *all-in-one*), a diffusion-based SBI method evaluated on four SBIBM benchmark tasks against four neural-SBI baselines (NLE, NPE, NPSE, NRE). It is a representative case: the method is equation-heavy, the evaluation requires rival baselines and generated datasets, and the reproduction passes through every stage of the pipeline. All figures below were obtained directly from two real snapshots of the same reproduction—one taken after synthesis (Stage 3) and one after repair (Stage 8.1)—using `diff -rq, find, and wc -l`.

Lifecycle of one component. The NLE baseline exhibits the canonical lifecycle. It is *identified* in the plan as one of four rival baselines (no code); the method’s own modules are *created* during synthesis while baselines do not yet exist; `src/baselines/nle.py` is then *created as a stub* delegating to a bundled fallback in the dataset-runner stage (Listing 1); the stub is *rewritten in place* into a faithful sbi-based implementation during repair, with the backup `nle.py.100.bak` retained; and finally a final-comparison test asserts the reproduced method outperforms NLE before the repository is frozen into the container task. The same trajectory holds for the dataset layout, paper constants, and missing algorithmic options—each *created* or *scaffolded* early and *corrected in place* at repair.

Quantitative signature of repair. Table 5 reports the Stage 3 $\rightarrow$ Stage 8.1 repository delta. The defining observation is that **the file count does not change and no net source files are added**: the repair stage is content-rewriting, not structure-adding, consistent with principle P4.

Representative rewrites. Listing 1 shows a baseline stub being rewritten into a faithful implementation during repair, and Listing 2 shows the dataset configuration being refactored from a single aggregated collection to per-task entries at paper scale.

Edited file	+	−	Nature of change
<code>configs/datasets.yaml</code>	129	53	collection $\rightarrow$ per-task; eval $\rightarrow$ test; paper-scale sizes
<code>src/diffusion/guidance.py</code>	35	9	implement optional self-recurrence refinement
<code>src/baselines/nle.py</code>	27	8	stub $\rightarrow$ faithful sbi NLE
<code>config.yaml</code>	1	0	add paper-mandated training_simulation

Table 6: Per-file deltas for the most significant edits. The largest change is configuration realignment, not new algorithms: structure is set early and repair *tunes* it.

Listing 1: `src/baselines/nle.py`: stub $\rightarrow$ faithful Neural Likelihood Estimation (+27/ $-$ 8).

```

# BEFORE (Stage 3): stub delegating to
bundled stack
def run_baseline(dataset, config=None):
    from src.baselines.
    simformer_posterior_only \
        import _run_repo_baseline
    overrides = dict(config or {})
    overrides["mask_mode"] = "posterior"
    return _run_repo_baseline(dataset=
dataset, config=overrides)

# AFTER (Stage 8.1): real NLE via the '
sbi' library
def run_baseline(dataset, config=None):
    # neural spline flow, batch size
    1000, Adam, early stopping
    seed = int((config or {}).get("seed",
0))
    batch_size = int((config or {}).get(
"batch_size", 1000))
    from src.baselines.sbi_baselines
    import run_sbi_nle
    return run_sbi_nle(dataset=dataset,
seed=seed, device=device,
batch_size=
batch_size, config=config or {})

```

Listing 2: `configs/datasets.yaml`: aggregated collection $\rightarrow$ per-task entries, sizes raised to paper scale (+129/ $-$ 53).

```

# BEFORE (Stage 3): single aggregated
collection
sbibm_tasks:
    kind: "task_collection"
    prep:
        splits: ["train", "eval"]
        train_num_samples_per_task: 10000
        eval_num_samples_per_task: 10
        tasks: ["gaussian_linear", "
gaussian_mixture",
"two_moons", "slcp"]

# AFTER (Stage 8.1): one entry per task;
eval -> test

```

```

1081 gaussian_linear:
1082     kind: "generated"
1083     simulator_task_name: "gaussian_linear"
1084     prep: { splits: ["train", "test"],
1085             train_num_samples: 20000,
1086             test_num_samples: 5000 }
1087     # gaussian_mixture / two_moons / slcp
1088     defined identically

```

These signatures—zero change in file count, a small net line delta dominated by a single config refactor, and one backup per edited file—confirm that, in this pipeline, behavioral fidelity is achieved by *rewriting a stable set of files*, not by regenerating the repository.

C Generated Test Examples

This section gives one representative generated unit test for each rubric category used in Figure 5. The examples are intentionally short: they illustrate the type of evidence encoded by each category, not the full generated test suite.

Environment, data, and assets. This BAM test checks that the JAX dependency is usable at a basic array-operation level.

```

1106 @pytest.mark.xfail(
1107     strict=False,
1108     reason="qualitative requirement;
1109     smoke test",
1110 )
1111 @pytest.mark.rubric(
1112     test_id="T-0001",
1113     rubric_id="env_jax",
1114     weight=1,
1115     tier="intermediate",
1116 )
1117 def test_env_jax():
1118     out = jnp.ones((2, 3)).shape
1119     assert tuple(out) == (2, 3)

```

Core method implementation. This BBox Adapter test checks a concrete invariant of the energy-based formulation: normalized probabilities sum to one.

```

1126 @pytest.mark.rubric(
1127     test_id="T-0012",
1128     rubric_id="impl_ebm_formulation",
1129     weight=1,
1130     tier="intermediate",
1131 )
1132 def test_impl_ebm_formulation():
1133     out = sum(ebm_normalized_probs(
1134         scores=[0.0, 0.0, 0.0],
1135         log_p_llm=[0.0, 0.0, 0.0]))
1136     assert out == 1.0

```

Baselines and metrics. This BAM test verifies that the metric entrypoints needed for paper-level comparisons are present and callable.

```

1142 @pytest.mark.xfail(
1143     strict=False,
1144     reason="qualitative requirement;
1145     smoke test",
1146 )
1147 @pytest.mark.rubric(
1148     test_id="T-0019",
1149     rubric_id="impl_metrics",
1150     weight=2,
1151     tier="intermediate",
1152 )
1153 def test_metrics_callable():
1154     for metric in [
1155         exact_gaussian_kls,
1156         compute_posterior_stats_metrics,
1157         reconstruction_mse,
1158     ]:
1159         assert callable(metric)

```

Experimental protocols and hyperparameters. This adaptive-pruning test directly checks that the small-model inference batch size matches the value specified by the paper.

```

1166 @pytest.mark.rubric(
1167     test_id="T-0020",
1168     rubric_id="hp_infer_bs_small_128",
1169     weight=1,
1170     tier="intermediate",
1171 )
1172 def test_infer_bs_small_128():
1173     out = EvalConfig().
1174         inference_batch_size_small
1175     assert out == 128.0

```

Experiment execution. This BBox test checks that the StrategyQA training and evaluation entrypoints exist, but it does not by itself prove that the full experiment has been run.

```

1182 @pytest.mark.xfail(
1183     strict=False,
1184     reason="qualitative requirement;
1185     smoke test",
1186 )
1187 @pytest.mark.rubric(
1188     test_id="T-0026",
1189     rubric_id="
1190     run_bbox_adapter_strategyqa",
1191     weight=1,
1192     tier="intermediate",
1193 )
1194 def test_bbox_strategyqa_entrypoints():
1195     assert callable(run_train)
1196     assert callable(run_eval)

```

Result reproduction and comparison. This BAM test corresponds to a reported figure, but only checks that the plotting entrypoint exists; stronger evidence would compare generated result files against the paper’s reported numbers or trends.




Figure 5: Distribution of generated PaperBench-style tests across six rubric categories for adaptive pruning, BAM, and BBox. Tests marked with `xfail` are counted as failures rather than being removed, while explicit `pytest.skip` cases are reported separately. The figure shows that the generated suites are dominated by expected-failing rubric checks, especially in method, asset, and result-reproduction categories, which supports our claim that execution-grounded paper-to-code evaluation must distinguish ordinary passing tests from qualitative smoke checks and incomplete reproduction tests.

```

1204 @pytest.mark.xfail(
1205     strict=False,
1206     reason="qualitative requirement;
1207     smoke test",
1208 )
1209
1210 @pytest.mark.rubric(
1211     test_id="T-0043",
1212     rubric_id="verify_fig_5_1",
1213     weight=1,
1214     tier="intermediate",
1215 )
1216 def test_fig_plotter_entrypoint():
1217     out = hasattr(
1218         Plotter,
1219         "plot_metric_trajectories",
1220     )
1221     assert out is True

```